

Problem Statement: Abalone Age Prediction

Description:-

Predicting the age of abalone from physical measurements. The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-consuming task. Other measurements, which are easier to obtain, are used to predict age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem.

Attribute Information:

Given is the attribute name, attribute type, measurement unit, and a brief description. The number of rings is the value to predict: either as a continuous value or as a classification problem.

Name / Data Type / Measurement Unit / Description

- 1- Sex / nominal / -- / M, F, and I (infant)
- 2- Length / continuous / mm / Longest shell measurement
- 3- Diameter / continuous / mm / perpendicular to length
- 4- Height / continuous / mm / with meat in shell
- 5- Whole weight / continuous / grams / whole abalone
- 6- Shucked weight / continuous / grams / weight of meat
- 7- Viscera weight / continuous / grams / gut weight (after bleeding)
- 8- Shell weight / continuous / grams / after being dried
- 9- Rings / integer / -- / +1.5 gives the age in years

Building a Regression Model

Import libraries

InÂ [1]:

```
import numpy as np

import pandas as pd

from scipy import stats

import matplotlib.pyplot as plt
```

```
import seaborn as sns

plt.style.use('seaborn-whitegrid')

%matplotlib inline
```

InÂ [2]:

```
from sklearn.preprocessing import OneHotEncoder, StandardScaler

from sklearn.decomposition import PCA

from sklearn.pipeline import Pipeline

from sklearn.compose import ColumnTransformer, make_column_selector

from sklearn.model_selection import train_test_split, cross_val_score,
GridSearchCV, RandomizedSearchCV
```

InÂ [3]:

```
from sklearn.neighbors import KNeighborsRegressor

from sklearn.linear_model import Ridge

from sklearn.svm import SVR

from sklearn.tree import DecisionTreeRegressor

from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor

from xgboost import XGBRegressor
```

InÂ [4]:

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

InÂ [5]:

```
import warnings

warnings.filterwarnings('ignore')
```

Download and load the dataset

InÂ [10]:

```
from google.colab import drive

drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).
```

In[12]:

```
data = pd.read_csv('drive/My Drive/abalone.csv')
```

```
data.head()
```

Out[12]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

In[13]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4177 entries, 0 to 4176
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Sex	4177 non-null	object
1	Length	4177 non-null	float64
2	Diameter	4177 non-null	float64
3	Height	4177 non-null	float64
4	Whole weight	4177 non-null	float64
5	Shucked weight	4177 non-null	float64
6	Viscera weight	4177 non-null	float64

```

7   Shell weight      4177 non-null   float64

8   Rings              4177 non-null   int64

dtypes: float64(7), int64(1), object(1)

memory usage: 293.8+ KB

```

Visualizations

Univariate Analysis

Univariate analysis provides an understanding in the characteristics of each feature in the data set. Different characteristics are computed for numerical and categorical data.

For the numerical features characteristics are standard deviation, skewness, kurtosis, percentile, interquartile range (IQR) and range.

In[14]:

```

stats_num = data.describe()

stats_num.loc['variance'] = data.select_dtypes(np.number).var().tolist()

stats_num.loc['skewness'] = data.select_dtypes(np.number).skew().tolist()

stats_num.loc['kurtosis'] = data.select_dtypes(np.number).kurtosis().tolist()

stats_num.loc['IQR'] = (data.select_dtypes(np.number).quantile(q=0.75) -
data.select_dtypes(np.number).quantile(q=0.25)).tolist()

stats_num.loc['range'] = (data.select_dtypes(np.number).max() -
data.select_dtypes(np.number).min()).tolist()

```

In[15]:

```
stats_num
```

Out[15]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000
variance	0.014422	0.009849	0.001750	0.240481	0.049268	0.012015	0.019377	10.395266
skewness	-0.639873	-0.609198	3.128817	0.530959	0.719098	0.591852	0.620927	1.114102
kurtosis	0.064621	-0.045476	76.025509	-0.023644	0.595124	0.084012	0.531926	2.330687
IQR	0.165000	0.130000	0.050000	0.711500	0.316000	0.159500	0.199000	3.000000
range	0.740000	0.595000	1.130000	2.823500	1.487000	0.759500	1.003500	28.000000

The feature Height has an example with value which doesn't make sense. This example needs to be removed.

InÂ [16]:

```
for column in data.select_dtypes(np.number).columns:
    p_value = stats.shapiro(data[column].dropna())[1]
    if p_value <= 0.05:
        print(f'Null hypothesis of normality for feature {column} is
rejected')
    else:
```

```
print(f'Null hypothesis of normality for feature {column} is  
accepted')
```

Null hypothesis of normality for feature Length is rejected

Null hypothesis of normality for feature Diameter is rejected

Null hypothesis of normality for feature Height is rejected

Null hypothesis of normality for feature Whole weight is rejected

Null hypothesis of normality for feature Shucked weight is rejected

Null hypothesis of normality for feature Viscera weight is rejected

Null hypothesis of normality for feature Shell weight is rejected

Null hypothesis of normality for feature Rings is rejected

The null hypothesis is rejected for every feature and the target, meaning that they aren't modelled with a normal distribution.

InÂ [17]:

```
data.hist(figsize=(15,5), layout=(2,4), bins=70)  
plt.show()
```

For the categorical features characteristics are count, cardinality, list of unique values, top and freq.

InÂ [18]:

```
stats_cat = data.select_dtypes('object').describe()
```

InÂ [19]:

```
stats_cat
```

Out[19]:

	Sex
count	4177
unique	3
top	M

Sex

freq 1528

InÂ [20]:

```
def uniqueValues(df):  
    for column in df:  
        unique_values = df[column].unique()  
        print(f'Unique values of feature {column} are: {unique_values}')
```

```
uniqueValues(data.select_dtypes('object'))
```

```
Unique values of feature Sex are: ['M' 'F' 'I']
```

InÂ [21]:

```
ax = sns.countplot(x=data['Sex'], data=data)  
  
for p in ax.patches:  
    x=p.get_bbox().get_points()[:,0]  
    y=p.get_bbox().get_points()[1,1]  
    ax.annotate(f'{p.get_height()}', (x.mean(), y), ha='center', va='bottom')  
  
plt.show()
```

Bi-Variate Analysis

InÂ [22]:

```
corr_matrix = data.corr()  
sns.heatmap(corr_matrix, annot=True)  
plt.show()
```

In this case it can be observed that there is somewhat weak linear relationship between each of the features and the target, so it is possible that linear models won't have a satisfactory performance.

Multi-Variate Analysis

InÂ [23]:

```
sns.pairplot(data, hue = 'Sex', diag_kind='hist')  
  
plt.show()
```

InÂ [24]:

```
nnum = data.select_dtypes(np.number).shape[1]  
  
cols = 4  
  
rows = 2  
  
fig, axes = plt.subplots(rows, cols, figsize=(20,10))  
  
for ax in axes.flatten():  
    ax.set_axis_off()  
  
for column, ax in zip(data.select_dtypes(np.number).columns, axes.flatten()):  
    sns.boxplot(x=data['Sex'], y=data[column], data=data, ax=ax)  
    ax.set_axis_on()  
  
plt.show()
```

Descriptive statistics

InÂ [26]:

```
data.describe()
```

Out[26]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
	0	0	0	0	0	0	0	0
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

Handle the Missing values

In[27]:

```
data.isnull().sum()
```

Out[27]:

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
```

```
Rings          0
```

```
dtype: int64
```

No missing values are found.

Find the outliers and replace the outliers

It was found that Height has a value of , so the whole data set is cleaned from examples that have a zero value because a physical measurment cannot have such value.

InÂ [28]:

```
data = data[(data != 0).all(axis=1)]
```

InÂ [29]:

```
data.describe()
```

Out[29]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4175.000000	4175.000000	4175.000000	4175.000000	4175.000000	4175.000000	4175.000000	4175.000000
	0	0	0	0	0	0	0	0
mean	0.524065	0.40794	0.139583	0.829005	0.359476	0.180653	0.238834	9.935090
std	0.120069	0.09922	0.041725	0.490349	0.221954	0.109605	0.139212	3.224227
min	0.075000	0.05500	0.010000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.35000	0.115000	0.442250	0.186250	0.093500	0.130000	8.000000
50%	0.545000	0.42500	0.140000	0.800000	0.336000	0.171000	0.234000	9.000000
75%	0.615000	0.48000	0.165000	1.153500	0.502000	0.253000	0.328750	11.000000
max	0.815000	0.65000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

Split the data into training and testing

InÂ [30]:

```
train, test = train_test_split(data, test_size=0.15, random_state=0)
```

InÂ [31]:

```
abaloneLen = len(data.index)
```

```
trainLen = len(train.index)
```

```
testLen = len(test.index)
```

```
trainPercent = np.round(trainLen/abaloneLen*100, 3)
```

```
print(trainPercent)
```

```
print(np.round(100 - trainPercent, 3))
```

```
print(abaloneLen, trainLen, testLen)
```

```
84.982
```

```
15.018
```

```
4175 3548 627
```

Split the data into dependent and independent variables

InÂ [32]:

```
X_train, y_train = train.drop('Rings', axis=1), train['Rings']
```

```
X_test, y_test = test.drop('Rings', axis=1), test['Rings']
```

Check for Categorical columns and perform encoding

InÂ [33]:

```
numerical_transformer = Pipeline([
    ('standardization', StandardScaler()),
    ('pca', PCA())
])
```

```
preprocessor = ColumnTransformer([
```

```

        ('lhot', OneHotEncoder(sparse=False),
make_column_selector(dtype_include='object')),

        ('num', numerical_transformer,
make_column_selector(dtype_include='number'))

])

```

InÂ [34]:

```

plt.subplot(1,2,1)

plt.hist(data['Rings'], bins=70)

plt.title('Target distribution')


plt.subplot(1,2,2)

plt.hist(np.log(data['Rings']), bins=70)

plt.title('Transformed target distribution')

plt.show()

```

InÂ [35]:

```

y_train_log, y_test_log = np.log(y_train), np.log(y_test)

```

Scale the independent variables

InÂ [36]:

```

check = pd.DataFrame(data=StandardScaler().fit_transform(X_train.drop('Sex',
axis=1)), columns=X_train.drop('Sex', axis=1).columns)

check['Rings'] = y_train_log


check.hist(figsize=(15,5), layout=(2,4), bins=70)

plt.show()

```

Build the Model

InÂ [37]:

```

from sklearn.model_selection import RepeatedKFold

from sklearn.compose import TransformedTargetRegressor

from sklearn.preprocessing import PowerTransformer


X, y = data.iloc[:, :-1], data.iloc[:, -1]


# minimally prepare dataset

y = y.astype('float32')


# evaluate model

baseline = SVR(kernel='rbf', gamma='scale', C=10)

transform = ColumnTransformer(transformers=[('c', OneHotEncoder(), [0])],
remainder='passthrough')

target = TransformedTargetRegressor(regressor=baseline,
transformer=PowerTransformer(), check_inverse=False)

pipeline = Pipeline(steps=[('ColumnTransformer', transform),
('Model', target)])

cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

r2_scores = cross_val_score(pipeline, X, y, scoring='r2', cv=cv, n_jobs=-1,
error_score='raise')

r2_scores = np.absolute(r2_scores)

print('Baseline: %.3f (+/-%.3f)' % (np.mean(r2_scores), np.std(r2_scores)))

Baseline: 0.562 (+/-0.033)

```

The regression algorithms that are trained to solve the problem are:

k-Nearest Neighbors;

Ridge Regression;

Support Vector Machines;

Decision Trees;

Random Forests;

Gradient Boosting for regression;

XGBoost Regressor.

10-fold cross-validation is used as a measure to prevent overfitting.

InÂ [38]:

```
models = [  
    KNeighborsRegressor(n_jobs=-1),  
    Ridge(),  
    SVR(),  
    DecisionTreeRegressor(random_state=0),  
    RandomForestRegressor(n_jobs=-1, random_state=0),  
    GradientBoostingRegressor(random_state=0),  
    XGBRegressor()  
]
```

Train the Model

InÂ [39]:

```
regressors = []  
  
for model in models:  
    regressor = Pipeline(steps=[('preprocessor', preprocessor), ('regressor',  
model)])  
    regressors.append(regressor)
```

InÂ [40]:

```
model_names = [  
    '    k-Neighbors Regressor',  
    '    Ridge Regression',  
    '    Support Vector Regressor',  
    '    Decision Tree Regressor',  
    '    Random Forest Regressor',
```

```

        'Gradient Boosting Regressor',
        'XGBoost Regressor'
    ]

for name, regressor in zip(model_names, regressors):

    regressor.fit(X_train, y_train_log)

    score = cross_val_score(estimator=regressor, X=X_train, y=y_train_log,
cv=10, scoring='r2')

    print(f"{name}: {score.mean():.4f} (+/-{score.std():.4f})")

    k-Neighbors Regressor: 0.5951 (+/-0.0425)

    Ridge Regression: 0.5866 (+/-0.0425)

    Support Vector Regressor: 0.6564 (+/-0.0359)

    Decision Tree Regressor: 0.3552 (+/-0.0959)

    Random Forest Regressor: 0.6545 (+/-0.0386)

    Gradient Boosting Regressor: 0.6628 (+/-0.0338)

[15:17:26] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.

[15:17:27] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.

[15:17:27] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.

[15:17:27] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.

[15:17:28] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.

[15:17:28] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.

[15:17:28] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.

[15:17:28] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.

```

```
[15:17:29] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.
```

```
[15:17:29] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.
```

```
[15:17:29] WARNING: /workspace/src/objective/regression_obj.cu:152:
reg:linear is now deprecated in favor of reg:squarederror.
```

```
XGBoost Regressor: 0.6629 (+/-0.0357)
```

Test the Model

InÂ [41]:

```
MAEs = []
```

```
MSEs = []
```

```
RMSEs = [] # If squared=True returns MSE value, if squared=False returns RMSE
value.
```

```
R2_scores = []
```

InÂ [42]:

```
for regressor in regressors:
```

```
    MAEs.append(mean_absolute_error(y_test_log, regressor.predict(X_test)))
```

```
    MSEs.append(mean_squared_error(y_test_log, regressor.predict(X_test)))
```

```
    RMSEs.append(mean_squared_error(y_test_log, regressor.predict(X_test),
squared=False))
```

```
    R2_scores.append(r2_score(y_test_log, regressor.predict(X_test)))
```

Measure the performance using metrics

InÂ [43]:

```
scoring_summary = pd.DataFrame({
    'Model': [name.strip() for name in model_names],
    'MAE score': MAEs,
    'MSE score': MSEs,
    'RMSE': RMSEs,
    'R2 Score': R2_scores
```



```
})
```

In[44]:

```
scoring_summary.sort_values('R2 Score', ascending=False)
```

Out[44]:

	Model	MAE score	MSE score	RMSE	R2 Score
2	Support Vector Regressor	0.141558	0.033549	0.183163	0.655027
6	XGBoost Regressor	0.140963	0.033856	0.183999	0.651870
5	Gradient Boosting Regressor	0.141429	0.034010	0.184418	0.650283
4	Random Forest Regressor	0.145248	0.035480	0.188362	0.635164
0	k-Neighbors Regressor	0.153654	0.040948	0.202356	0.578944
1	Ridge Regression	0.159779	0.042259	0.205570	0.565463
3	Decision Tree Regressor	0.191218	0.064029	0.253039	0.341607