

Problem Statement: Customer Segmentation Analysis

Problem Statement:

You own the mall and want to understand the customers who can quickly converge [Target Customers] so that the insight can be given to the marketing team and plan the strategy accordingly.

Import libraries

InÂ [1]:

```
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
```

Load the dataset

InÂ [2]:

```
from google.colab import drive
drive.mount('/content/drive')
Mounted at /content/drive
```

InÂ [3]:

```
data = pd.read_csv('drive/My Drive/Mall_Customers.csv')
```

```
data.head()
```

Out[3]:

CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
------------	--------	-----	---------------------	------------------------

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

In[4]:

```
data.shape
```

Out[4]:

```
(200, 5)
```

In[5]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 200 entries, 0 to 199
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	CustomerID	200 non-null	int64
1	Gender	200 non-null	object
2	Age	200 non-null	int64
3	Annual Income (k\$)	200 non-null	int64
4	Spending Score (1-100)	200 non-null	int64

```
dtypes: int64(4), object(1)
```

```
memory usage: 7.9+ KB
```

Visualizations

Univariate Analysis

Univariate analysis provides an understanding in the characteristics of each feature in the data set. Different characteristics are computed for numerical and categorical data.

For the numerical features characteristics are standard deviation, skewness, kurtosis, percentile, interquartile range (IQR) and range.

InÂ [6]:

```
stats_num = data.describe()

stats_num.loc['variance'] =
data.select_dtypes(np.number).var().tolist()

stats_num.loc['skewness'] =
data.select_dtypes(np.number).skew().tolist()

stats_num.loc['kurtosis'] =
data.select_dtypes(np.number).kurtosis().tolist()

stats_num.loc['IQR'] =
(data.select_dtypes(np.number).quantile(q=0.75) -
data.select_dtypes(np.number).quantile(q=0.25)).tolist()

stats_num.loc['range'] = (data.select_dtypes(np.number).max() -
data.select_dtypes(np.number).min()).tolist()
```

InÂ [7]:

```
stats_num
```

Out[7]:

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000
variance	3350.000000	195.133166	689.835578	666.854271
skewness	0.000000	0.485569	0.321843	-0.047220
kurtosis	-1.200000	-0.671573	-0.098487	-0.826629
IQR	99.500000	20.250000	36.500000	38.250000
range	199.000000	52.000000	122.000000	98.000000

InÂ [8]:

```
from scipy import stats

for column in data.select_dtypes(np.number).columns:
    p_value = stats.shapiro(data[column].dropna()) [1]
    if p_value <= 0.05:
```

```
        print(f'Null hypothesis of normality for feature  
{column} is rejected')
```

```
    else:
```

```
        print(f'Null hypothesis of normality for feature  
{column} is accepted')
```

Null hypothesis of normality for feature CustomerID is rejected

Null hypothesis of normality for feature Age is rejected

Null hypothesis of normality for feature Annual Income (k\$) is rejected

Null hypothesis of normality for feature Spending Score (1-100) is rejected

The null hypothesis is rejected for every feature and the target, meaning that they aren't modelled with a normal distribution.

InÂ [9]:

```
data.hist(figsize=(15,5), layout=(2,4), bins=70)
```

```
plt.show()
```

For the categorical features characteristics are count, cardinality, list of unique values, top and freq.

InÂ [10]:

```
stats_cat = data.select_dtypes('object').describe()
```

InÂ [11]:

```
stats_cat
```

Out[11]:

Gender

Gender

count 200

unique 2

top Female

freq 112

InÂ [12]:

```
def uniqueValues(df):  
    for column in df:  
        unique_values = df[column].unique()  
        print(f'Unique values of feature {column} are:  
{unique_values}')
```

```
uniqueValues(data.select_dtypes('object'))
```

```
Unique values of feature Gender are: ['Male' 'Female']
```

InÂ [13]:

```
ax = sns.countplot(x=data['Gender'], data=data)
```

```
for p in ax.patches:
```

```
    x=p.get_bbox().get_points()[0,0]
```

```
    y=p.get_bbox().get_points()[1,1]
```

```
    ax.annotate(f'{p.get_height()}', (x.mean(), y), ha='center',  
va='bottom')
```

```
plt.show()
```

Bi-variate Analysis

InÂ [14]:

```
corr_matrix = data.corr()  
sns.heatmap(corr_matrix, annot=True)  
plt.show()
```

Multi-variate Analysis

InÂ [18]:

```
sns.pairplot(data, hue = 'Gender', diag_kind='hist')  
plt.show()
```

InÂ [19]:

```
nnum = data.select_dtypes(np.number).shape[1]  
cols = 2  
rows = 2  
fig, axes = plt.subplots(rows, cols, figsize=(20,10))  
  
for ax in axes.flatten():  
    ax.set_axis_off()  
  
for column, ax in zip(data.select_dtypes(np.number).columns,  
axes.flatten()):
```

```
sns.boxplot(x=data['Gender'], y=data[column], data=data,  
ax=ax)
```

```
ax.set_axis_on()
```

```
plt.show()
```

Descriptive statistics

InÂ [17]:

```
data.describe()
```

Out[17]:

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

Check for missing values

InÂ [20]:


```
data.isnull().sum()
```

Out[20]:

```
CustomerID          0
Gender              0
Age                0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```

No missing values are found.

Finding outliers and scaling data

InÂ [21]:

```
x = data['CustomerID']
y = data['Annual Income (k$)']
```

```
plt.plot(x, y)
```

Out[21]:

```
[<matplotlib.lines.Line2D at 0x7f657391aa10>]
```

InÂ [22]:

```
x = data['Annual Income (k$)']
y = data['Age']
```

```
plt.plot(x, y)
```

Out[22]:

```
[<matplotlib.lines.Line2D at 0x7f657388d910>]
```

Split the data into dependent and independent variables

InÂ [23]:

```
x = data.iloc[:, [3, 4]].values
```

```
print(x.shape)
```

```
(200, 2)
```

Clustering algorithms

KMeans clustering

InÂ [24]:

```
from sklearn.cluster import KMeans
```

```
wcss = []
```

```
for i in range(1, 11):
```

```
    km = KMeans(n_clusters = i, init = 'k-means++', max_iter =  
300, n_init = 10, random_state = 0)
```

```
    km.fit(x)
```

```
    wcss.append(km.inertia_)
```

```
plt.plot(range(1, 11), wcss)
```

```
plt.title('The Elbow Method')
```

```
plt.xlabel('No. of Clusters')
```

```
plt.ylabel('wcss')
```

```
plt.show()
```

InÂ [25]:

```
km = KMeans(n_clusters = 5, init = 'k-means++', max_iter = 300,
n_init = 10, random_state = 0)

y_means = km.fit_predict(x)

plt.scatter(x[y_means == 0, 0], x[y_means == 0, 1], s = 100, c =
'pink', label = 'miser')

plt.scatter(x[y_means == 1, 0], x[y_means == 1, 1], s = 100, c =
'yellow', label = 'general')

plt.scatter(x[y_means == 2, 0], x[y_means == 2, 1], s = 100, c =
'cyan', label = 'target')

plt.scatter(x[y_means == 3, 0], x[y_means == 3, 1], s = 100, c =
'magenta', label = 'spendthrift')

plt.scatter(x[y_means == 4, 0], x[y_means == 4, 1], s = 100, c =
'orange', label = 'careful')

plt.scatter(km.cluster_centers_[0,0], km.cluster_centers_[0, 1],
s = 50, c = 'blue' , label = 'centeroid')

plt.title('K Means Clustering')

plt.xlabel('Annual Income')

plt.ylabel('Spending Score')

plt.legend()

plt.show()
```

Hierarchical clustering

InÂ [26]:

```
import scipy.cluster.hierarchy as sch
```

```

dendrogram = sch.dendrogram(sch.linkage(x, method = 'ward'))
plt.title('Dendrogram')
plt.xlabel('Customers')
plt.ylabel('Euclidean Distance')
plt.show()

```

InÂ [27]:

```

from sklearn.cluster import AgglomerativeClustering

hc = AgglomerativeClustering(n_clusters = 5, affinity =
'euclidean', linkage = 'ward')

y_hc = hc.fit_predict(x)

plt.scatter(x[y_means == 0, 0], x[y_means == 0, 1], s = 100, c =
'pink', label = 'miser')

plt.scatter(x[y_means == 1, 0], x[y_means == 1, 1], s = 100, c =
'yellow', label = 'general')

plt.scatter(x[y_means == 2, 0], x[y_means == 2, 1], s = 100, c =
'cyan', label = 'target')

plt.scatter(x[y_means == 3, 0], x[y_means == 3, 1], s = 100, c =
'magenta', label = 'spendthrift')

plt.scatter(x[y_means == 4, 0], x[y_means == 4, 1], s = 100, c =
'orange', label = 'careful')

plt.scatter(km.cluster_centers_[:,0], km.cluster_centers_[:, 1],
s = 50, c = 'blue' , label = 'centroid')

plt.title('Hierarchial Clustering')
plt.xlabel('Annual Income')

```

```
plt.ylabel('Spending Score')
plt.legend()
plt.show()
```

Clusters of customers based on their ages

Build the Model

InÂ [28]:

```
data.columns
```

Out[28]:

```
Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',
      'Spending Score (1-100)'],
      dtype='object')
```

InÂ [29]:

```
x = data.iloc[:, [2, 4]].values
x.shape
```

Out[29]:

```
(200, 2)
```

Train the Model

InÂ [30]:

```
from sklearn.cluster import KMeans
```

```
wcss = []
```

```
for i in range(1, 11):
```

```
    kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter =
300, n_init = 10, random_state = 0)
```

```
kmeans.fit(x)

wcss.append(kmeans.inertia_)


plt.rcParams['figure.figsize'] = (7, 5)
plt.plot(range(1, 11), wcss)
plt.title('K-Means Clustering(The Elbow Method)', fontsize = 20)
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```

Test the Model

InÂ [31]:

```
kmeans = KMeans(n_clusters = 4, init = 'k-means++', max_iter =
300, n_init = 10, random_state = 0)

ymeans = kmeans.fit_predict(x)


plt.rcParams['figure.figsize'] = (10, 10)
plt.title('Cluster of Ages', fontsize = 30)


plt.scatter(x[ymeans == 0, 0], x[ymeans == 0, 1], s = 100, c =
'pink', label = 'Usual Customers' )

plt.scatter(x[ymeans == 1, 0], x[ymeans == 1, 1], s = 100, c =
'orange', label = 'Priority Customers')

plt.scatter(x[ymeans == 2, 0], x[ymeans == 2, 1], s = 100, c =
'lightgreen', label = 'Target Customers(Young)')

plt.scatter(x[ymeans == 3, 0], x[ymeans == 3, 1], s = 100, c =
'red', label = 'Target Customers(Old)')
```

```
plt.scatter(kmeans.cluster_centers_[ :, 0],  
kmeans.cluster_centers_[ :, 1], s = 50, c = 'black')
```

```
plt.xlabel('Age')  
plt.ylabel('Spending Score (1-100)')  
plt.legend()  
plt.show()
```

Clusters of customers based on their gender

Check for Categorical columns and perform encoding.

InÂ [32]:

```
data['Gender'].replace(['Male', 'Female'], [0, 1], inplace =  
True)
```

```
data['Gender'].value_counts()
```

Out[32]:

```
1      112
```

```
0       88
```

```
Name: Gender, dtype: int64
```

Split the data into dependent and independent variables.

InÂ [33]:

```
x = data.iloc[:, [1, 4]].values
```

```
x.shape
```

Out[33]:

```
(200, 2)
```

Train the Model

InÂ [34]:

```
from sklearn.cluster import KMeans

wcss = []

for i in range(1, 11):

    kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter =
300, n_init = 10, random_state = 0)

    kmeans.fit(x)

    wcss.append(kmeans.inertia_)


plt.rcParams['figure.figsize'] = (7, 7)
plt.title('The Elbow Method', fontsize = 20)
plt.plot(range(1, 11), wcss)
plt.xlabel('No. of Clusters', fontsize = 10)
plt.ylabel('wcss')
plt.show()
```

Test the Model

InÂ [35]:

```
kmeans = KMeans(n_clusters = 3, max_iter = 300, n_init = 10,
random_state = 0)

ymeans = kmeans.fit_predict(x)


plt.rcParams['figure.figsize'] = (10, 10)

plt.scatter(x[ymeans == 0, 0], x[ymeans == 0, 1], s = 80, c =
'pink', label = 'low spending score')
```



```
plt.scatter(x[ymeans == 1, 0], x[ymeans == 1, 1], s = 80, c =
'orange', label = 'medium spending score')

plt.scatter(x[ymeans == 2, 0], x[ymeans == 2, 1], s = 80, c =
'lightgreen', label = 'high spending score')

plt.scatter(kmeans.cluster_centers_[0,0],
kmeans.cluster_centers_[0,1], s = 50, color = 'blue')

plt.legend()

plt.show()
```

From the above cluster plot, we can conclude that males and females belong to all categories (high, low and medium spending score category).