

## Data Visualization and Pre-processing

### Import libraries

InÂ [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

### Load dataset

InÂ [2]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

InÂ [3]:

```
data = pd.read_csv('drive/My Drive/Churn_Modelling.csv')
```

```
data.head()
```

Out[3]:

	Row Number	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	157013	Bonini	699	France	Female	39	1	0.00	2	0	0	93826.63	0

	Row Number	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
		54				ale								
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

InÂ [4]:

```
data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender                10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                10000 non-null  int64
8   Balance               10000 non-null  float64
9   NumOfProducts         10000 non-null  int64
10  HasCrCard             10000 non-null  int64
11  IsActiveMember        10000 non-null  int64
12  EstimatedSalary       10000 non-null  float64
13  Exited                10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

## Visualisations

### 1. Univariate Analysis

InÂ [5]:

```
data['Gender'].value_counts()
```

Out[5]:

```
Male      5457
Female    4543
```

Name: Gender, dtype: int64

InÂ [7]:

```
# Plotting the features of the dataset to see the correlation  
between them
```

```
plt.hist(x = data.Gender, bins = 3, color = 'pink')  
plt.title('comparison of male and female')  
plt.xlabel('Gender')  
plt.ylabel('population')  
plt.show()
```

InÂ [6]:

```
data['Age'].value_counts()
```

Out[6]:

```
37      478  
38      477  
35      474  
36      456  
34      447
```

```
...  
92         2  
82         1  
88         1  
85         1  
83         1
```

Name: Age, Length: 70, dtype: int64

InÂ [8]:

```
# comparison of age in the dataset
```

```
plt.hist(x = data.Age, bins = 10, color = 'orange')  
plt.title('comparison of Age')  
plt.xlabel('Age')  
plt.ylabel('population')  
plt.show()
```

InÂ [9]:

```
data['Geography'].value_counts()
```

Out[9]:

```
France      5014  
Germany     2509  
Spain       2477
```

Name: Geography, dtype: int64

InÂ [10]:

```
# comparison of geography
```

```
plt.hist(x = data.Geography, bins = 5, color = 'green')
plt.title('comparison of Geography')
plt.xlabel('Geography')
plt.ylabel('population')
plt.show()
```

InÂ [11]:

```
data['HasCrCard'].value_counts()
```

Out[11]:

```
1    7055
0    2945
Name: HasCrCard, dtype: int64
```

InÂ [12]:

```
# comparision of how many customers hold the credit card

plt.hist(x = data.HasCrCard, bins = 3, color = 'red')
plt.title('how many people have or not have the credit card')
plt.xlabel('customers holding credit card')
plt.ylabel('population')
plt.show()
```

InÂ [13]:

```
data['IsActiveMember'].value_counts()
```

Out[13]:

```
1    5151
0    4849
Name: IsActiveMember, dtype: int64
```

InÂ [14]:

```
# How many active member does the bank have ?

plt.hist(x = data.IsActiveMember, bins = 3, color = 'brown')
plt.title('Active Members')
plt.xlabel('Customers')
plt.ylabel('population')
plt.show()
```

## 2. Bi - Variate Analysis

InÂ [15]:

```
# comparison between Geography and Gender
```

```
Gender = pd.crosstab(data['Gender'],data['Geography'])
```

```
Gender.div(Gender.sum(1).astype(float), axis=0).plot(kind="bar",
stacked=True, figsize=(6, 6))
```

Out[15]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a93dbbfd0>
```

In [16]:

```
# comparison between geography and card holders
```

```
HasCrCard = pd.crosstab(data['HasCrCard'], data['Geography'])
HasCrCard.div(HasCrCard.sum(1).astype(float), axis =
0).plot(kind = 'bar',
stacked =
True,figsize = (6, 6))
```

Out[16]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a93ced590>
```

In [17]:

```
# comparison of active member in differnt geographies
```

```
IsActiveMember = pd.crosstab(data['IsActiveMember'],
data['Geography'])
IsActiveMember.div(IsActiveMember.sum(1).astype(float), axis =
0).plot(kind = 'bar',
stacked = True,
figsize= (6, 6))
```

Out[17]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a93c7c950>
```

In [18]:

```
# comparing ages in different geographies
```

```
Age = pd.crosstab(data['Age'], data['Geography'])
Age.div(Age.sum(1).astype(float), axis = 0).plot(kind = 'bar',
stacked = True,
figsize = (15,15))
```

Out[18]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a93bfea10>
```

In [19]:

```
# calculating total balance in france, germany and spain
```

```
total_france = data.Balance[data.Geography == 'France'].sum()
total_germany = data.Balance[data.Geography == 'Germany'].sum()
total_spain = data.Balance[data.Geography == 'Spain'].sum()
```

```
print("Total Balance in France :",total_france)
print("Total Balance in Germany :",total_germany)
print("Total Balance in Spain :",total_spain)

Total Balance in France : 311332479.49
Total Balance in Germany : 300402861.38
Total Balance in Spain : 153123552.01
```

InÂ [20]:

```
# plotting a pie chart

labels = 'France', 'Germany', 'Spain'
colors = ['cyan', 'magenta', 'orange']
sizes = [311, 300, 153]
explode = [ 0.01, 0.01, 0.01]

plt.pie(sizes, colors = colors, labels = labels, explode =
explode, shadow = True)

plt.axis('equal')
plt.show()
```

### 3. Multi - Variate Analysis

InÂ [21]:

```
sns.pairplot(data=data, hue='Exited')
```

Out[21]:

```
<seaborn.axisgrid.PairGrid at 0x7f6a93ddd510>
```

### Descriptive statistics

InÂ [23]:

```
#Statistical analysis
data.describe()
```

Out[23]:

	Row Num ber	Cust omer Id	Cred itSco re	Age	Tenu re	Balan ce	Num OfPro ducts	Has CrC ard	IsActi veMe mber	Estim atedS alary	Exite d
<b>co</b>	1000	1.000	1000	1000	1000	10000	10000.	1000	10000.	10000.	1000
<b>u</b>	0.00	000e	0.000	0.000	0.000	.0000	00000	0.00	00000	00000	0.000
<b>nt</b>	000	+04	000	000	000	00	0	000	0	0	000

	Row Num ber	Cust omer Id	Cred itSco re	Age	Tenu re	Balan ce	Num OfPro ducts	Has CrC ard	IsActi veMe mber	Estim atedS alary	Exite d
m ea n	5000 .500 00	1.569 094e +07	650.5 2880 0	38.92 1800	5.012 800	76485 .8892 88	1.5302 00	0.70 550	0.5151 00	10009 0.2398 81	0.203 700
st d	2886 .895 68	7.193 619e +04	96.65 3299	10.48 7806	2.892 174	62397 .4052 02	0.5816 54	0.45 584	0.4997 97	57510. 49281 8	0.402 769
m in	1.00 000	1.556 570e +07	350.0 0000 0	18.00 0000	0.000 000	0.000 000	1.0000 00	0.00 000	0.0000 00	11.580 000	0.000 000
2 5 %	2500 .750 00	1.562 853e +07	584.0 0000 0	32.00 0000	3.000 000	0.000 000	1.0000 00	0.00 000	0.0000 00	51002. 11000 0	0.000 000
5 0 %	5000 .500 00	1.569 074e +07	652.0 0000 0	37.00 0000	5.000 000	97198 .5400 00	1.0000 00	1.00 000	1.0000 00	10019 3.9150 00	0.000 000
7 5 %	7500 .250 00	1.575 323e +07	718.0 0000 0	44.00 0000	7.000 000	12764 4.240 000	2.0000 00	1.00 000	1.0000 00	14938 8.2475 00	0.000 000
m a x	1000 0.00 000	1.581 569e +07	850.0 0000 0	92.00 0000	10.00 0000	25089 8.090 000	4.0000 00	1.00 000	1.0000 00	19999 2.4800 00	1.000 000

## Handle the Missing values

```
#Missing Values
data.isnull().sum()
```

InÂ [24]:

Out[24]:

RowNumber 0

```
CustomerId      0
Surname         0
CreditScore     0
Geography      0
Gender          0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited         0
dtype: int64
No missing values are found.
```

### Find the outliers and replace the outliers

```
sns.boxplot(data = data, x = 'CreditScore')
```

InÂ [25]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a8ecfe7d0>
```

Out[25]:

```
sns.boxplot(data = data, x = 'Age')
```

InÂ [26]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a8e9a3f50>
```

Out[26]:

```
sns.boxplot(data = data, x = 'Balance')
```

InÂ [27]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a8d0e0e90>
```

Out[27]:

```
sns.boxplot(data = data, x = 'EstimatedSalary')
```

InÂ [28]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a8d0e0c50>
```

Out[28]:

```
for i in data:
    if data[i].dtype=='int64' or data[i].dtypes=='float64':
        q1=data[i].quantile(0.25)
        q3=data[i].quantile(0.75)
```

InÂ [29]:



```

iqr=q3-q1
upper=q3+1.5*iqr
lower=q1-1.5*iqr
data[i]=np.where(data[i] >upper, upper, data[i])
data[i]=np.where(data[i] <lower, lower, data[i])

```

In [30]:

```
data.describe()
```

Out[30]:

	Row Number	Customer Id	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	1000	1000	1000	1000	1000	10000	10000	1000	10000	10000	1000
	0.000	0.000e+04	0.000000	0.000000	0.000000	.000000	000000	0.000000	000000	000000	0.000000
mean	5000.5000	1.569094e+07	650.561300	38.660800	5.012800	76485.889288	1.527200	0.70550	0.515100	100090.239881	0.0
std	2886.89568	7.193619e+04	96.558702	9.746704	2.892174	62397.405202	0.570081	0.45584	0.499797	57510.492818	0.0
min	1.00000	1.556570e+07	383.000000	18.000000	0.000000	0.000000	1.000000	0.000000	0.000000	11.580000	0.0
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000	0.000000	51002.110000	0.0
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.000000	1.000000	100193.915000	0.0
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.000000	1.000000	149388.247500	0.0

	Row Num ber	Cust omer Id	Cred itSco re	Age	Tenu re	Balan ce	NumO fProd ucts	Has CrC ard	IsActi veMe mber	Estim atedSa lary	Ex ite d
m	1000	1.581	850.0			25089				19999	
a	0.000	569e	0000	62.00	10.00	8.090	3.5000	1.00	1.0000	2.4800	0.0
x	00	+07	0	0000	0000	000	00	000	00	00	

## Preprocessing

InÂ [31]:

```
# Removing the unnecassary features from the dataset

data = data.drop(['CustomerId', 'Surname', 'RowNumber'], axis =
1)

print(data.columns)
Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure',
'Balance',
      'NumOfProducts', 'HasCrCard', 'IsActiveMember',
'EstimatedSalary',
      'Exited'],
      dtype='object')
```

InÂ [32]:

```
data.shape
```

Out[32]:

```
(10000, 11)
```

## Split the data into dependent and independent variables

InÂ [33]:

```
# splitting the dataset into x(independent variables) and
y(dependent variables)

x = data.iloc[:,0:10]
y = data.iloc[:,10]

print(x.shape)
print(y.shape)

print(x.columns)
(10000, 10)
(10000,)
```

```
Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure',
      'Balance',
      'NumOfProducts', 'HasCrCard', 'IsActiveMember',
      'EstimatedSalary'],
      dtype='object')
```

### Check for Categorical columns and perform encoding

InÂ [34]:

```
# Encoding Categorical variables into numerical variables
# One Hot Encoding
```

```
x = pd.get_dummies(x)
```

```
x.head()
```

Out[34]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Geography_France	Geography_Germany	Geography_Spain	Gender_Female	Gender_Male
0	619.0	42.0	2.0	0.00	1.0	1.0	1.0	101348.88	1	0	0	1	0
1	608.0	41.0	1.0	83807.86	1.0	0.0	1.0	112542.58	0	0	1	1	0
2	502.0	42.0	8.0	159660.80	3.0	1.0	0.0	113931.57	1	0	0	1	0
3	699.0	39.0	1.0	0.00	2.0	0.0	0.0	93826.63	1	0	0	1	0

	Credit Score	Age	Te en ur e	Ba lan ce	Num OfPr oduc ts	Ha sCr Ca rd	IsAct iveM embe r	Esti mate dSal ary	Geogr aphy_ Franc e	Geogr aphy_ Germa ny	Geog raph y_Sp ain	Gen der_ Fem ale	Gen der_ _M ale
4	850	43	2	1255	1.0	1.0	1.0	7908	0	0	1	1	0
	.0	.	0	10.				4.10					
		0		82									

## Split the data into training and testing

InÂ [35]:

```
# splitting the data into training and testing set

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size = 0.25, random_state = 0)

print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

(7500, 13)
(7500,)
(2500, 13)
(2500,)
```

## Scale the independent variables

InÂ [36]:

```
# Feature Scaling
# Only on Independent Variable to convert them into values
ranging from -1 to +1

from sklearn.preprocessing import StandardScaler

sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.fit_transform(x_test)

x_train = pd.DataFrame(x_train)
x_train.head()
```

												Out[36]:	
	0	1	2	3	4	5	6	7	8	9	10	11	12
<b>0</b>	- 0.73 682 8	- 0.04 228 3	- 0.00 886 0	- 0.67 316 0	- 2.58 323 1	- 1.55 362 4	- 1.03 446 0	- 1.64 081 0	- 1.01 558 8	- 1.76 021 6	- 0.57 468 2	- 1.08 726 1	- 1.08 726 1
<b>1</b>	- 1.02 525 7	- 0.67 449 6	- 0.00 886 0	- 1.20 772 4	- 0.82 257 8	- 0.64 365 7	- 1.03 446 0	- 0.07 927 2	- 0.98 465 1	- 0.56 811 2	- 0.57 468 2	- 1.08 726 1	- 1.08 726 1
<b>2</b>	- 0.80 886 1	- 0.46 970 2	- 1.39 329 3	- 0.35 693 7	- 0.82 257 8	- 0.64 365 7	- 0.96 668 8	- 0.99 684 0	- 1.01 558 8	- 0.56 811 2	- 1.74 009 4	- 1.08 726 1	- 1.08 726 1
<b>3</b>	- 0.39 667 7	- 0.06 011 4	- 0.00 886 0	- 0.00 935 6	- 0.93 807 6	- 0.64 365 7	- 0.96 668 8	- 1.59 174 6	- 1.01 558 8	- 0.56 811 2	- 1.74 009 4	- 0.91 974 3	- 0.91 974 3
<b>4</b>	- 0.46 890 8	- 1.37 344 4	- 0.70 107 7	- 1.20 772 4	- 0.82 257 8	- 0.64 365 7	- 0.96 668 8	- 1.28 330 2	- 0.98 465 1	- 0.56 811 2	- 0.57 468 2	- 0.91 974 3	- 0.91 974 3